

Lending Club Loan Data Analysis

July 21, 2023

Beacon Nicolas Mkhabele Project

```
[3]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[4]: data = pd.read_csv("loan_data.csv")
```

```
[5]: data.head()
```

```
[5]:   credit.policy      purpose  int.rate  installment  log.annual.inc \
0           1  debt_consolidation    0.1189     829.10    11.350407
1           1        credit_card    0.1071     228.22    11.082143
2           1  debt_consolidation    0.1357     366.86    10.373491
3           1  debt_consolidation    0.1008     162.34    11.350407
4           1        credit_card    0.1426     102.92    11.299732

      dti  fico  days.with.cr.line  revol.bal  revol.util  inq.last.6mths \
0  19.48    737       5639.958333     28854      52.1                 0
1  14.29    707       2760.000000     33623      76.7                 0
2  11.63    682       4710.000000     3511      25.6                 1
3   8.10    712       2699.958333     33667      73.2                 1
4  14.97    667       4066.000000     4740      39.5                 0

      delinq.2yrs  pub.rec  not.fully.paid
0            0        0                  0
1            0        0                  0
2            0        0                  0
3            0        0                  0
4            1        0                  0
```

```
[6]: #understanding the dataset
data.shape
```

```
[6]: (9578, 14)
```

```
[7]: data.isnull().sum()
```

```
[7]: credit.policy      0
purpose          0
int.rate          0
installment       0
log.annual.inc    0
dti              0
fico             0
days.with.cr.line 0
revol.bal         0
revol.util        0
inq.last.6mths    0
delinq.2yrs       0
pub.rec           0
not.fully.paid    0
dtype: int64
```

```
[8]: #Transforming categorical values into numerical values (discrete)
data.sample(10)
```

```
[8]:   credit.policy      purpose  int.rate  installment \
4157          1 debt_consolidation  0.1426    514.59
908           1 debt_consolidation  0.1197    292.16
1899          1 debt_consolidation  0.0963    320.95
3951          1 credit_card        0.1568    700.04
5580          1 credit_card        0.1565    131.20
3137          1 home_improvement  0.1284    53.79
5084          1 all_other          0.1218    333.00
4552          1 debt_consolidation  0.1287    504.50
4680          1 debt_consolidation  0.1218    666.00
6094          1 debt_consolidation  0.1114    518.30

      log.annual.inc    dti  fico  days.with.cr.line  revol.bal  revol.util \
4157     10.858999  19.68  687      3313.000000    22918     90.9
908      10.819778   8.14  697      6840.000000    8243     37.0
1899     11.156251   6.31  737      3480.041667   20423     46.7
3951     11.695247  16.77  677      2280.000000   29799     53.0
5580      9.952278  18.91  697      1170.000000   13881     90.7
3137     10.561008   2.46  682      3180.000000    2076     83.0
5084     11.144814  18.22  712      5820.000000   17604     72.1
4552     11.652687   8.35  702      6089.958333   12152     54.2
4680     10.836950  22.88  732      3990.000000   26803     51.4
6094     11.512925  21.88  747      4620.000000   29860     47.6


      inq.last.6mths  delinq.2yrs  pub.rec  not.fully.paid
4157            1            0          0            0
908            2            0          1            0
1899            2            0          0            1
```

```
3951      3      0      0      0
5580      0      0      0      0
3137      1      0      0      0
5084      0      0      0      0
4552      2      0      1      0
4680      1      0      0      0
6094      0      0      0      0
```

```
[9]: np.unique(data['credit.policy'], return_counts=True)
```

```
[9]: (array([0, 1]), array([1868, 7710]))
```

```
[10]: plt.figure(figsize=(7, 4))
sns.countplot(x=data['credit.policy']);
```



```
[11]: data.dtypes
```

```
[11]: credit.policy        int64
purpose            object
int.rate         float64
installment       float64
log.annual.inc    float64
dti              float64
fico             int64
days.with.cr.line float64
revol.bal        int64
```

```
revol.util           float64
inq.last.6mths      int64
delinq.2yrs          int64
pub.rec              int64
not.fully.paid       int64
dtype: object
```

```
[12]: df_col_info = pd.DataFrame(data.dtypes, columns= ['col_data_type'])

df_col_info_names_changed = df_col_info.reset_index()

df_col_info_names_changed.columns = ['col_names', 'col_data_type']
df_col_info_names_changed
```


```
[12]:      col_names col_data_type
0        credit.policy      int64
1            purpose        object
2            int.rate     float64
3        installment     float64
4    log.annual.inc     float64
5                dti     float64
6                fico     int64
7  days.with.cr.line   float64
8        revol.bal      int64
9        revol.util     float64
10       inq.last.6mths   int64
11       delinq.2yrs      int64
12         pub.rec      int64
13    not.fully.paid     int64
```

```
[13]: df_col_info_count = pd.DataFrame(df_col_info_names_changed.
                                         ↪groupby("col_data_type").count().reset_index())

df_col_info_count.columns= ['col_name', 'dtype']
df_col_info_count
```

```
[13]:  col_name  dtype
0      int64    7
1    float64    6
2      object    1
```

```
[14]: plt.figure(figsize=(17, 9))
sns.countplot(y=data['purpose']);
```



```
[15]: #Null value Analysis
NAs= pd.concat([data.isnull().sum()], axis=1)

NAs[NAs.sum(axis=1)>0]
```

[15]: Empty DataFrame
Columns: [0]
Index: []

```
[16]: pd.set_option('display.float_format', lambda x:'%.4f' % x)
```


```
[17]: data.describe().T
```

	count	mean	std	min	25%	\
credit.policy	9578.0000	0.8050	0.3962	0.0000	1.0000	
int.rate	9578.0000	0.1226	0.0268	0.0600	0.1039	
installment	9578.0000	319.0894	207.0713	15.6700	163.7700	
log.annual.inc	9578.0000	10.9321	0.6148	7.5475	10.5584	
dti	9578.0000	12.6067	6.8840	0.0000	7.2125	
fico	9578.0000	710.8463	37.9705	612.0000	682.0000	
days.with.cr.line	9578.0000	4560.7672	2496.9304	178.9583	2820.0000	
revol.bal	9578.0000	16913.9639	33756.1896	0.0000	3187.0000	
revol.util	9578.0000	46.7992	29.0144	0.0000	22.6000	
inq.last.6mths	9578.0000	1.5775	2.2002	0.0000	0.0000	
delinq.2yrs	9578.0000	0.1637	0.5462	0.0000	0.0000	
pub.rec	9578.0000	0.0621	0.2621	0.0000	0.0000	
not.fully.paid	9578.0000	0.1601	0.3667	0.0000	0.0000	
		50%	75%		max	

```
credit.policy      1.0000    1.0000    1.0000
int.rate          0.1221    0.1407    0.2164
installment       268.9500  432.7625  940.1400
log.annual.inc   10.9289   11.2913   14.5284
dti              12.6650   17.9500   29.9600
fico             707.0000  737.0000  827.0000
days.with.cr.line 4139.9583 5730.0000 17639.9583
revol.bal        8596.0000 18249.5000 1207359.0000
revol.util       46.3000   70.9000   119.0000
inq.last.6mths   1.0000   2.0000   33.0000
delinq.2yrs      0.0000   0.0000   13.0000
pub.rec          0.0000   0.0000   5.0000
not.fully.paid   0.0000   0.0000   1.0000
```


```
[18]: data.installment.hist()
```

```
[18]: <AxesSubplot:>
```




```
[19]: data.dti.hist()
```

```
[19]: <AxesSubplot:>
```



```
[20]: data['pub.rec'].hist()
```

```
[20]: <AxesSubplot:>
```



```
[21]: # evaluating the usefulness of the records
```

```
data.dtypes
```

```
[21]: credit.policy      int64
purpose            object
int.rate        float64
installment     float64
log.annual.inc   float64
dti              float64
fico             int64
days.with.cr.line float64
revol.bal        int64
revol.util       float64
inq.last.6mths    int64
delinq.2yrs      int64
pub.rec           int64
not.fully.paid    int64
dtype: object
```

```
[22]: data.columns
```

```
[22]: Index(['credit.policy', 'purpose', 'int.rate', 'installment', 'log.annual.inc',
       'dti', 'fico', 'days.with.cr.line', 'revol.bal', 'revol.util',
       'inq.last.6mths', 'delinq.2yrs', 'pub.rec', 'not.fully.paid'],
      dtype='object')
```

```
[23]: col_to_evaluate = ['credit.policy', 'purpose', 'int.rate', 'installment', 'log.
       ↴annual.inc',
       'dti', 'fico', 'days.with.cr.line', 'revol.bal', 'revol.util']
```

```
[24]: from scipy import stats
```

```
[25]: # Forming 2 datasets from the original dataset as subsets
```

```
class1 = data[(data['credit.policy'] == 0)]
class1.reset_index(inplace= True)

class2 = data[(data['credit.policy'] == 1)]
class2.reset_index(inplace = True)
```

```
[26]: stats.ttest_ind(class1['int.rate'],class2['int.rate'])[1]
```

```
[26]: 2.1945619717925775e-190
```

Since the predicted value is much lower than 5% then it is useful for Machine learning model

this means we can create a model for the columns

```
[27]: for each_col in col_to_evaluate:  
    col1 = pd.to_numeric(class1[each_col], errors='coerce')  
    col2 = pd.to_numeric(class2[each_col], errors='coerce')  
  
    p_value = stats.ttest_ind(col1, col2, nan_policy='omit')[1]  
  
    if p_value < 0.05:  
        # Column is useful  
        print("The column '{}' is useful with a p-value of {}".format(each_col,  
        ↪p_value))  
    else:  
        # Column is not useful  
        print("The column '{}' is not useful with a p-value of {}".format(each_col, p_value))
```

```
The column 'credit.policy' is useful with a p-value of 0.0  
The column 'purpose' is not useful with a p-value of 0.0  
The column 'int.rate' is useful with a p-value of 2.1945619717925775e-190  
The column 'installment' is useful with a p-value of 8.620919919293774e-09  
The column 'log.annual.inc' is useful with a p-value of 0.0006337324172009995  
The column 'dti' is useful with a p-value of 4.945272027510337e-19  
The column 'fico' is useful with a p-value of 2.6100416830751396e-271  
The column 'days.with.cr.line' is useful with a p-value of 2.627126143480234e-22  
The column 'revol.bal' is useful with a p-value of 1.592869483344647e-76  
The column 'revol.util' is useful with a p-value of 1.7179159037327666e-24  
  
/usr/local/lib/python3.7/site-packages/scipy/stats/mstats_basic.py:1050:  
RuntimeWarning: divide by zero encountered in true_divide  
    denom = ma.sqrt(svar*(1.0/n1 + 1.0/n2)) # n-D computation here!
```

```
[28]: cols_to_evaluate2 = ['inq.last.6mths', 'delinq.2yrs', 'pub.rec', 'not.fully.  
    ↪paid']
```

Categorical usefulness

```
[29]: from scipy.stats import chi2_contingency
```

```
[30]: #Tests for the relationships
```

```
chi_res = chi2_contingency(pd.crosstab(data['credit.policy'], data['not.fully.  
    ↪paid']))
```

```
[31]: chi_res
```

```
[31]: (238.3788010698609,
       8.87573133930704e-54,
       1,
       array([[1569.01858426, 298.98141574],
              [6475.98141574, 1234.01858426]]))

[32]: chi2_contingency(pd.crosstab(data['credit.policy'], data['not.fully.paid']))[1]

[32]: 8.87573133930704e-54

[33]: for each_col in cols_to_evaluate2 :

    chi2_contingency(pd.crosstab(data['credit.policy'], data['not.fully.
    ↪paid']))[1]

    if p_value < 0.05:
        # Column is useful
        print("The column '{}' is useful with a p-value of {}".format(each_col, ↪
    ↪p_value))
    else:
        # Column is not useful
        print("The column '{}' is not useful with a p-value of {}".format(
    ↪each_col, p_value))
```

The column 'inq.last.6mths' is useful with a p-value of 1.7179159037327666e-24
 The column 'delinq.2yrs' is useful with a p-value of 1.7179159037327666e-24
 The column 'pub.rec' is useful with a p-value of 1.7179159037327666e-24
 The column 'not.fully.paid' is useful with a p-value of 1.7179159037327666e-24

encoding cat columns

```
[34]: data_dummies= pd.get_dummies(data)

[35]: data_dummies.sample(10)

[35]:   credit.policy  int.rate  installment  log.annual.inc  dti  fico \
8206          0     0.1425      102.9000      10.9519  22.1700   657
935           1     0.1008      161.5300      10.2036  12.7100   732
4740          1     0.1287      571.7700      11.5229  17.2200   712
7096          1     0.1311      404.9400      10.7013  22.4300   692
1615           1     0.1367      204.1100      10.7096  22.7800   692
419            1     0.0775      305.9700      11.0021  9.2000   797
3106          1     0.1095      327.1400      10.4009  24.8800   727
8284          0     0.1229      80.0500       9.1695  0.0000   662
4494          1     0.1357      203.8200      10.9681  9.3700   692
8671          0     0.1387      218.3400      10.6454  17.3700   657
```

	days.with.cr.line	revol.bal	revol.util	inq.last.6mths	delinq.2yrs	\
8206	4949.0417	6710	54.6000	1	0	
935	3150.0000	11837	20.6000	0	0	
4740	7979.9583	26949	52.0000	0	0	
7096	5490.0000	10337	61.9000	0	0	
1615	8490.0417	19270	97.3000	1	0	
419	6270.0000	209	1.9000	0	0	
3106	3726.9583	10462	33.9000	1	0	
8284	539.0417	0	0.0000	5	0	
4494	1410.0000	902	20.0000	1	0	
8671	1710.0000	8312	85.7000	0	1	

	pub.rec	not.fully.paid	purpose_all_other	purpose_credit_card	\
8206	0	0	1	0	
935	0	0	0	0	
4740	0	0	0	0	
7096	0	0	0	0	
1615	1	0	0	1	
419	0	0	0	1	
3106	0	0	0	0	
8284	0	0	1	0	
4494	0	0	1	0	
8671	0	0	0	1	

	purpose_debt_consolidation	purpose_educational	\
8206	0	0	
935	1	0	
4740	1	0	
7096	1	0	
1615	0	0	
419	0	0	
3106	1	0	
8284	0	0	
4494	0	0	
8671	0	0	

	purpose_home_improvement	purpose_major_purchase	purpose_small_business
8206	0	0	0
935	0	0	0
4740	0	0	0
7096	0	0	0
1615	0	0	0
419	0	0	0
3106	0	0	0
8284	0	0	0
4494	0	0	0
8671	0	0	0

```
[36]: data_dummies.sample(10).T
```

	3394	7269	6585	347	\
credit.policy	1.0000	1.0000	1.0000	1.0000	
int.rate	0.1126	0.1348	0.1183	0.1141	
installment	328.6400	393.5200	795.2200	16.4700	
log.annual.inc	11.2252	11.0186	11.7753	9.8782	
dti	20.0600	10.9000	15.3800	3.6900	
fico	717.0000	682.0000	812.0000	667.0000	
days.with.cr.line	4529.9583	4560.0417	11705.0000	8820.0000	
revol.bal	3404.0000	3715.0000	346.0000	12229.0000	
revol.util	72.4000	26.5000	0.7000	90.6000	
inq.last.6mths	2.0000	1.0000	0.0000	0.0000	
delinq.2yrs	0.0000	0.0000	0.0000	0.0000	
pub.rec	0.0000	0.0000	0.0000	1.0000	
not.fully.paid	0.0000	0.0000	0.0000	0.0000	
purpose_all_other	0.0000	0.0000	0.0000	1.0000	
purpose_credit_card	0.0000	1.0000	0.0000	0.0000	
purpose_debt_consolidation	1.0000	0.0000	1.0000	0.0000	
purpose_educational	0.0000	0.0000	0.0000	0.0000	
purpose_home_improvement	0.0000	0.0000	0.0000	0.0000	
purpose_major_purchase	0.0000	0.0000	0.0000	0.0000	
purpose_small_business	0.0000	0.0000	0.0000	0.0000	
	4837	7371	5518	3755	\
credit.policy	1.0000	1.0000	1.0000	1.0000	
int.rate	0.1218	0.1533	0.1322	0.1189	
installment	599.4000	870.7100	507.0100	33.1700	
log.annual.inc	11.6082	11.2772	10.3810	11.3737	
dti	8.2300	13.9600	15.6300	5.5700	
fico	722.0000	697.0000	707.0000	697.0000	
days.with.cr.line	5429.9583	9721.0417	14879.9583	8519.9583	
revol.bal	15533.0000	26018.0000	14489.0000	5742.0000	
revol.util	79.3000	43.9000	89.4000	55.2000	
inq.last.6mths	1.0000	0.0000	0.0000	0.0000	
delinq.2yrs	0.0000	1.0000	0.0000	1.0000	
pub.rec	0.0000	0.0000	0.0000	0.0000	
not.fully.paid	0.0000	0.0000	0.0000	0.0000	
purpose_all_other	0.0000	0.0000	0.0000	1.0000	
purpose_credit_card	0.0000	0.0000	0.0000	0.0000	
purpose_debt_consolidation	1.0000	1.0000	1.0000	0.0000	
purpose_educational	0.0000	0.0000	0.0000	0.0000	
purpose_home_improvement	0.0000	0.0000	0.0000	0.0000	
purpose_major_purchase	0.0000	0.0000	0.0000	0.0000	
purpose_small_business	0.0000	0.0000	0.0000	0.0000	

credit.policy	1.0000	1.0000
int.rate	0.1197	0.0743
installment	199.2000	124.3000
log.annual.inc	10.9819	11.0429
dti	6.0000	5.8200
fico	672.0000	772.0000
days.with.cr.line	4410.0000	5430.0000
revol.bal	3831.0000	773.0000
revol.util	34.2000	4.8000
inq.last.6mths	0.0000	0.0000
delinq.2yrs	1.0000	0.0000
pub.rec	0.0000	0.0000
not.fully.paid	1.0000	0.0000
purpose_all_other	0.0000	0.0000
purpose_credit_card	0.0000	0.0000
purpose_debt_consolidation	1.0000	0.0000
purpose_educational	0.0000	0.0000
purpose_home_improvement	0.0000	1.0000
purpose_major_purchase	0.0000	0.0000
purpose_small_business	0.0000	0.0000

```
[37]: from sklearn.preprocessing import StandardScaler
      from sklearn.compose import ColumnTransformer
```

```
[38]: x_df = data.drop(['credit.policy'], axis=1)
      y_df = data['credit.policy']
```

Building a Model

```
[39]: data_transf = ColumnTransformer([('std-sclr', StandardScaler(), ['credit.policy'],
                                     ['purpose', 'int.rate', 'installment', 'log.annual.inc',
                                      'dti', 'fico', 'days.with.cr.line', 'revol.bal', 'revol.util']),
                                     remainder = 'passthrough')]
```

0.0.1 Scaling

```
[40]: x_df_encoded = pd.get_dummies(x_df, columns=['purpose'])

      data_transf = StandardScaler()
```

```
[41]: X_scaled = data_transf.fit_transform(x_df_encoded)
```

```
[42]: y = y_df.values
```

Using multi-collinearity

```
[43]: y = data['credit.policy'].values

[44]: from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn.model_selection import train_test_split
from sklearn import metrics
from sklearn.ensemble import RandomForestClassifier

[45]: X = data.drop(['credit.policy','purpose'],axis = 1)

[46]: X.values

[46]: array([[1.18900000e-01, 8.29100000e+02, 1.13504065e+01, ...,
       0.00000000e+00, 0.00000000e+00, 0.00000000e+00],
      [1.07100000e-01, 2.28220000e+02, 1.10821426e+01, ...,
       0.00000000e+00, 0.00000000e+00, 0.00000000e+00],
      [1.35700000e-01, 3.66860000e+02, 1.03734912e+01, ...,
       0.00000000e+00, 0.00000000e+00, 0.00000000e+00],
      ...,
      [1.07100000e-01, 9.78100000e+01, 1.05966347e+01, ...,
       0.00000000e+00, 0.00000000e+00, 1.00000000e+00],
      [1.60000000e-01, 3.51580000e+02, 1.08197783e+01, ...,
       0.00000000e+00, 0.00000000e+00, 1.00000000e+00],
      [1.39200000e-01, 8.53430000e+02, 1.12644641e+01, ...,
       0.00000000e+00, 0.00000000e+00, 1.00000000e+00]])
```

```
[47]: mod = pd.DataFrame()
mod['VIF_Factor'] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
mod['feature'] = X.columns
mod = mod.sort_values(by='VIF_Factor', ascending=False).round(5)

print(mod)
```

	VIF_Factor	feature
2	377.0961	log.annual.inc
4	278.9056	fico
0	35.2293	int.rate
7	5.6752	revol.util
5	5.2466	days.with.cr.line
3	5.0872	dti
1	4.2039	installment
8	1.6652	inq.last.6mths
6	1.5606	revol.bal
11	1.2471	not.fully.paid
9	1.2000	delinq.2yrs
10	1.1004	pub.rec

```
[48]: X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size= .25)
```

```
[49]: ml_model = RandomForestClassifier()

ml_model.fit(X_train,y_train)

y_pred_train = ml_model.predict(X_train)
y_pred = ml_model.predict(X_test)

print("train_ ing accuracy score:", metrics.accuracy_score(y_train, y_pred_train))
print("train_ ing accuracy score:", metrics.accuracy_score(y_test, y_pred))

train_ ing accuracy score: 1.0
train_ ing accuracy score: 0.9870563674321503
```

```
[50]: print(metrics.classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	1.00	0.94	0.97	457
1	0.99	1.00	0.99	1938
accuracy			0.99	2395
macro avg	0.99	0.97	0.98	2395
weighted avg	0.99	0.99	0.99	2395

```
[51]: metrics.confusion_matrix(y_test, y_pred)
```

```
[51]: array([[ 428,    29],
           [    2, 1936]])
```

Random Foresting the model

```
[52]: num_trees_list = [1,2,3,4,5,7,10]
```

```
[53]: train_ = []
test_ = []

for num_trees in num_trees_list:
    tres_= RandomForestClassifier(n_estimators = num_trees)

    tres_.fit(X_train,y_train)

    train_pred = tres_.predict(X_train)
```

```
order = metrics.accuracy_score(y_train, train_pred)

train_.append(order)
test_pred = tres_.predict(X_test)

order = metrics.accuracy_score(y_test, test_pred)

test_.append(order)
```

```
[54]: num_estimators = [3,5,7,9,11,13,15,17,19,21]
max_depths = [2,4,6,8,10,12,14]
min_samples_split = np.linspace(.01,.05, 6)
```

```
[55]: parameters = dict(
    max_depths = max_depths,
    min_samples_split = min_samples_split,
    estimatead = num_estimators
)
print(parameters)
```

```
{'max_depths': [2, 4, 6, 8, 10, 12, 14], 'min_samples_split': array([0.01 , 0.018, 0.026, 0.034, 0.042, 0.05 ]), 'estimatead': [3, 5, 7, 9, 11, 13, 15, 17, 19, 21]}
```

```
[56]: from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestClassifier
```

```
[57]: model = RandomForestClassifier()
```

```
[58]: parameters = {
    'n_estimators': [10, 50, 100],
    'max_depth': [None, 5, 10],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
}
```

```
Grid = GridSearchCV(
    estimator=model,
    param_grid=parameters,
    scoring='neg_mean_squared_error',
    cv=5,
    verbose=1
)
Grid.fit(X_scaled, y)
```

```
Fitting 5 folds for each of 81 candidates, totalling 405 fits
```

```
[58]: GridSearchCV(cv=5, estimator=RandomForestClassifier(),
                  param_grid={'max_depth': [None, 5, 10],
                               'min_samples_leaf': [1, 2, 4],
                               'min_samples_split': [2, 5, 10],
                               'n_estimators': [10, 50, 100]},
                  scoring='neg_mean_squared_error', verbose=1)
```

```
[59]: Grid.best_params_
```

```
[59]: {'max_depth': None,
       'min_samples_leaf': 1,
       'min_samples_split': 5,
       'n_estimators': 100}
```

```
[60]: ml_model = RandomForestClassifier(
          max_depth=None,
          min_samples_leaf=1,
          min_samples_split=5,
          n_estimators=50
          #Balanced
      )

ml_model.fit(X_train,y_train)

y_pred_train = ml_model.predict(X_train)
y_pred = ml_model.predict(X_test)

print("train ing accuracy score:", metrics.accuracy_score(y_train, y_pred_train))
print("train ing accuracy score:", metrics.accuracy_score(y_test, y_pred))
```

```
train ing accuracy score: 0.9979117360434359
```

```
train ing accuracy score: 0.9870563674321503
```

```
[61]: metrics.confusion_matrix(y_test, y_pred)
```

```
[61]: array([[ 429,    28],
           [    3, 1935]])
```

```
[66]: from sklearn.linear_model import LogisticRegression
```

```
[67]: ml_model = LogisticRegression()
```

```

ml_model.fit(X_train,y_train)

y_pred_train = ml_model.predict(X_train)
y_pred = ml_model.predict(X_test)

print("train ing accuracy score:", metrics.accuracy_score(y_train, y_pred))
print("train ing accuracy score:", metrics.accuracy_score(y_test, y_pred))

```

train ing accuracy score: 0.9050535987748851
train ing accuracy score: 0.9006263048016702

Building a model using deep learning

[68]:

```

from keras.models import Sequential
from keras.layers import Dense

```

[69]:

```
X_train.shape
```

[69]:

```
(7183, 19)
```

[70]:

```

model = Sequential()

model.add(Dense(512, input_dim=19, activation='relu'))
model.add(Dense(256, activation='relu'))
model.add(Dense(128, activation='relu'))
model.add(Dense(64, activation='relu'))
model.add(Dense(32, activation='relu'))
model.add(Dense(16, activation='relu'))
model.add(Dense(8, activation='relu'))
model.add(Dense(1, activation='sigmoid'))

model.compile(loss='binary_crossentropy',
              optimizer='adam',
              metrics=['accuracy'])

```

[71]:

```
from keras.callbacks import EarlyStopping
```

[72]:

```
E_stop = EarlyStopping(monitor='val_loss', mode ='min', patience= 5, verbose= 1)
```

[73]:

```
%%time
```

```

history = model.fit(
    X_train,
    y_train,
    epochs= 100,
    batch_size=25,

```

```

        validation_data= (X_test,y_test),
        verbose=1,
        callbacks=[E_stop]
)

Epoch 1/100
288/288 [=====] - 7s 5ms/step - loss: 0.2868 -
accuracy: 0.8881 - val_loss: 0.2806 - val_accuracy: 0.9027
Epoch 2/100
288/288 [=====] - 1s 4ms/step - loss: 0.2224 -
accuracy: 0.9120 - val_loss: 0.2127 - val_accuracy: 0.9161
Epoch 3/100
288/288 [=====] - 1s 4ms/step - loss: 0.1892 -
accuracy: 0.9247 - val_loss: 0.1953 - val_accuracy: 0.9223
Epoch 4/100
288/288 [=====] - 1s 4ms/step - loss: 0.1690 -
accuracy: 0.9367 - val_loss: 0.1522 - val_accuracy: 0.9424
Epoch 5/100
288/288 [=====] - 1s 4ms/step - loss: 0.1482 -
accuracy: 0.9415 - val_loss: 0.1393 - val_accuracy: 0.9549
Epoch 6/100
288/288 [=====] - 1s 4ms/step - loss: 0.1279 -
accuracy: 0.9522 - val_loss: 0.1464 - val_accuracy: 0.9407
Epoch 7/100
288/288 [=====] - 1s 4ms/step - loss: 0.1159 -
accuracy: 0.9564 - val_loss: 0.1216 - val_accuracy: 0.9541
Epoch 8/100
288/288 [=====] - 1s 4ms/step - loss: 0.1099 -
accuracy: 0.9589 - val_loss: 0.1168 - val_accuracy: 0.9570
Epoch 9/100
288/288 [=====] - 1s 4ms/step - loss: 0.0991 -
accuracy: 0.9635 - val_loss: 0.1190 - val_accuracy: 0.9557
Epoch 10/100
288/288 [=====] - 1s 4ms/step - loss: 0.0948 -
accuracy: 0.9667 - val_loss: 0.1060 - val_accuracy: 0.9599
Epoch 11/100
288/288 [=====] - 1s 4ms/step - loss: 0.0857 -
accuracy: 0.9719 - val_loss: 0.1090 - val_accuracy: 0.9608
Epoch 12/100
288/288 [=====] - 1s 4ms/step - loss: 0.0778 -
accuracy: 0.9710 - val_loss: 0.1154 - val_accuracy: 0.9516
Epoch 13/100
288/288 [=====] - 1s 4ms/step - loss: 0.0767 -
accuracy: 0.9720 - val_loss: 0.1017 - val_accuracy: 0.9670
Epoch 14/100
288/288 [=====] - 1s 4ms/step - loss: 0.0618 -
accuracy: 0.9765 - val_loss: 0.1019 - val_accuracy: 0.9587

```

```
Epoch 15/100
288/288 [=====] - 1s 4ms/step - loss: 0.0604 -
accuracy: 0.9758 - val_loss: 0.1218 - val_accuracy: 0.9608
Epoch 16/100
288/288 [=====] - 1s 4ms/step - loss: 0.0548 -
accuracy: 0.9806 - val_loss: 0.1331 - val_accuracy: 0.9495
Epoch 17/100
288/288 [=====] - 1s 4ms/step - loss: 0.0484 -
accuracy: 0.9811 - val_loss: 0.1273 - val_accuracy: 0.9612
Epoch 18/100
288/288 [=====] - 1s 4ms/step - loss: 0.0541 -
accuracy: 0.9812 - val_loss: 0.1231 - val_accuracy: 0.9653
Epoch 18: early stopping
CPU times: user 35.1 s, sys: 4.9 s, total: 40 s
Wall time: 27.9 s
```

```
[74]: plt.style.use('ggplot')


def plot_history(history):
    acc = history.history['accuracy']
    val_acc = history.history['val_accuracy']
    loss = history.history['loss']
    val_loss = history.history['val_loss']

    x = range(1, len(acc) + 1)

    plt.figure(figsize=(11, 7))
    plt.subplot(1, 2, 1)
    plt.plot(x, acc, 'b', label='Training Accuracy')
    plt.plot(x, val_acc, 'r', label='Validation Accuracy')
    plt.title('Training and Validation Accuracy')
    plt.legend()

    plt.subplot(1, 2, 2)
    plt.plot(x, loss, 'b', label='Training Loss')
    plt.plot(x, val_loss, 'r', label='Validation Loss')
    plt.title('Training and Validation Loss')
    plt.legend()
```

```
[75]: plot_history(history)
```



```
[76]: y_pred = np.squeeze((y_pred > .5 ).astype('int32'))
```

```
[77]: from sklearn import metrics
```

```
[78]: metrics.accuracy_score(y_test, y_pred)
```

```
[78]: 0.9006263048016702
```

```
[ ]:
```